

Wide but Shallow

Technical-Interview Problems Are Four Ideas in Many Costumes

Conor Garvey Gelvin*
gelvin@outlook.com

Data snapshot: June 2026

Abstract

Interview preparation is taught as a long catalogue of separate problem “patterns”—sliding window, monotonic stack, prefix sums, and dozens more—each memorized on its own. We show the catalogue is mostly costume: those “patterns” are really **four underlying ideas**: a single left-to-right pass (a *fold*), divide-and-solve recursion (dynamic programming), spreading information across a graph until it settles, and binary search. To put a number on this without hand-waving, we take a *pre-registered* random sample of 100 problems drawn from seven major technology firms’ most-asked lists (769 distinct problems) and, for each, prove with the Lean proof assistant that a standard accepted solution really is one of the four ideas. **71 of the 100 carry such a proof** (69 machine-checked here, 2 citing existing formal proofs); the other 29 are a genuine “tail” that fits no single idea. Every classification was checked against the problem’s published editorial. So about **three-quarters of interview problems reduce to four ideas**, each backed by a machine-checked proof, and the single pass is the most common—many tricks taught as unrelated (prefix sums, the XOR trick, a hash “seen set,” reversing a list, the sliding window) are the *same* pass, differing only in what running value they carry. A secondary observation: in our pooled lists the seven firms ask almost the same mix (bias-corrected Cramér’s V of 0.091, above the 0.065 that sampling noise alone would produce but far below “moderate”; only Uber stands out)—though this similarity may partly be an artifact of how the platform aggregates and truncates the lists, rather than evidence about any firm’s true private bar. We describe which problems are *commonly asked*, not how hard they are or whether company-specific drilling pays off. In short: interview problems are *wide on the surface but shallow underneath*. All derived data and proofs are released; the raw scrape is withheld per the platform’s terms.

1 Two myths

Two beliefs dominate interview preparation. The first, and our main target, is *pattern proliferation*: that the field is a long list of loosely related tricks—sliding window, monotonic stack, union-find, prefix sums, and so on—each learned separately, as popular pattern catalogues present them [1]. The second is *company specificity*: that Google, Meta, Amazon and their peers test different skills, so you should prepare differently for each. Both are widely repeated and, as far as we can tell, never tested.

We test them, and then ask what underlying *structure* explains the answer. The short version: the “dozens of patterns” you are told to memorize are a handful of ideas in different disguises;

*The author interviewed at three of the firms studied, passing all three, and came to these four ideas by noticing them recur while practising—not by applying them by design.

and, at least in our pooled data, the seven companies ask nearly the same things. To keep our labels honest we do not hand-sort the problems and stop there—we use a proof assistant to check, problem by problem, that the idea we assign really does solve the problem (and, for most of them, that the solution is fully correct).

2 Data and method

For seven firms (Amazon, Apple, Bloomberg, Google, Meta, Microsoft, Uber) we scraped a public interview-prep platform (LeetCode) for how often each problem appears among that firm’s *most-asked* problems (the platform exposes a top-200 list per firm and recency window, so rarely-asked problems are truncated; we note this where it matters). Across firms this covers 769 distinct problems, broken down by recency. The platform tags each one with a *surface family* (“sliding window,” “monotonic stack,” “dynamic programming,” and so on), which we consolidate into about twenty broad families.

We map each problem to one of four underlying *ideas* (formally, *recursion schemes*) by the *shape* of its standard solution, or to a leftover “tail” for problems that fit no single idea (Table 1). We ask two separate questions and keep them apart. The first is *coverage*: what fraction of problems do the four ideas explain? This is settled problem by problem with proofs, not labels. For each problem in a pre-registered random sample of 100, we attempt to prove that a standard accepted solution really is one of the four ideas (Section 4); coverage is the fraction of proofs that succeed. A rule written in advance (Appendix A) picks only which idea to attempt first — it never decides whether a problem counts. The second question is the *company* comparison: do firms ask different mixes? For this we count how often each firm asks problems of each idea — a seven-by-four table with $n = 1,400$ asked-problem entries — and measure how different the seven rows are (Section 3). The comparison is repeated over the finer twenty families as a robustness check, and every difference measure is bias-corrected [7] and reported with uncertainty.

3 Result: four ideas, not dozens

Wide surface. The surface taxonomy really is spread out: no one or two families dominate (it takes six of the ~ 20 families to cover even half of asked problems; counting each distinct problem once instead, the spread is similarly flat—Gini 0.30, where 0 is a perfectly even split and 1 is total concentration). By the usual measure, interview preparation *looks* like a long, flat list. This is the breadth the prep industry sells.

Shallow root. But that breadth is mostly repackaging. Sorting the same problems by the *shape* of their standard solution, about three-quarters reduce to just four ideas—**71 of a random sample of 100, each backed by a machine-checked proof** (Section 4; 95% Wilson interval [61%, 79%])—with the single left-to-right pass the most common. Tricks memorized as unrelated—prefix sums, the XOR trick, a hash “seen set,” reversing a list, matching parentheses with a stack, the sliding window—are the *same* pass; they differ only in what running value they carry. The variety is in the bookkeeping, not in the computation. The remaining quarter—bespoke data structures, simulation, geometry, ad-hoc number theory—is a genuine tail we do not force into a scheme.

One shared syllabus. The same data dissolve the company myth. We score how differently the firms ask on a standard 0-to-1 scale (bias-corrected Cramér’s V [6, 7], where 0 means identical and the usual cutoffs call 0.1 “small” and 0.3 “moderate”). Across the four ideas the firms score

Underlying idea	in scheme	formal test (Lean)	textbook families it subsumes
streaming fold	27	IsFold	two-pointers, sliding-window, monotonic-stack, hashing, prefix-sum, bit-XOR, string-scan
recursive decomposition (DP)	25	catamorphism	dynamic programming, backtracking, tree, trie
graph relaxation	11	least fixpoint	BFS/DFS, Dijkstra, Bellman–Ford, topo-sort, union–find
bisection	8	monotone threshold	binary search (monotone predicate)
one of the four	71		
not a scheme	—		design, simulation, heaps, number theory, geometry, string matching, ...

Table 1: The roughly twenty surface families collapse onto four recursion schemes. On a pre-registered sample of 100 problems, 71 are proven to be one of the four (69 machine-checked here, 2 cited [14, 16]; both graph relaxation); 29 form the genuine tail. “Formal test” names the Lean predicate each certificate proves membership in. Full method and numbers: Sections 3–4.

just $V = 0.091$ (95% bootstrap interval $[0.06, 0.12]$)—small, with the whole interval well below “moderate,” so the difference is genuinely bounded, not merely undetected. We report the four-idea grouping because its 0.091 stands clear of its 0.065 chance level (the score sampling noise alone would produce); the finer twenty-family score falls *below* its own, larger chance level (0.064 vs. 0.116)—indistinguishable from identical—so the finer view, if anything, makes the firms look more alike. Comparing firms two at a time, the only real differences all involve **Uber**, which asks more graph problems and fewer single-pass ones (its largest gap, versus Google, is still only “small”; Figure 1); **the other six firms are statistically indistinguishable**. One caveat: we compare the *mix of problem types*, not how hard the problems are or whether drilling for a specific company helps—our data cannot measure those. Within that scope, there is essentially one interview syllabus, and preparing differently per company is, for most firms, wasted effort.

4 Machine-checking the classification

A classification is only as good as its labels, and the usual check—a second human rater—only tells you whether two people read a label the same way. We do something stronger: we settle each label with a *proof assistant*, a tool that mechanically verifies mathematical claims. Working in Lean 4 [8] with its library Mathlib [9], we write each of the four ideas as a precise test on programs (for example, “this function is a single left-to-right pass”). Then, for each sampled problem, we model a standard accepted solution and prove two things: it has the *shape* of its idea, and it satisfies a property of the *specific* problem—its actual grid, array, or condition. The second part gives the proof teeth: an abstract statement would certify the *category*, not the problem. So every certificate is held to a mechanical standard: a released Lean program (**lake exe gate**) recomputes each verdict from the compiled proofs — both theorems must be about the modelled solution itself; a kernel-checked




Figure 1: How much each firm over- or under-asks each idea, relative to the shared average (standardized residuals; red = more than expected, blue = less). Only Uber stands out—it asks far more graph problems (+5.7) and far fewer single-pass ones (−3.4)—while the other six rows are muted (all within about ± 2). This is the picture behind the small overall difference ($V = 0.091$): one shared syllabus, with Uber the lone exception.

evaluation of that solution on a concrete input (usually the problem’s published example) must be present, which no abstract placeholder can supply; and only Lean’s standard axioms may be used, ruling out unfinished proofs and compiler-trusted shortcuts. No human judgement decides whether a finished certificate counts.

Could “streaming fold” be so loose that everything qualifies? No, and that is itself a theorem: a function is a fold exactly when its answer on a longer input is determined by its answer so far plus the one new element (`isFold_iff_update`), and this test can be failed — carrying only a running count provably cannot count distinct elements (`not_isFold_countDistinct`). We never demand a proof that the algorithm is *optimal*: the platform’s acceptance vouches for the submissions themselves, and our question is only *which idea solves the problem* — though most certificates prove exact correctness anyway (49 of 69, below).

A pre-registered sample. Which problems land in the sample matters, because some are far easier to formalize than others. An earlier draft used a sample we had already spent weeks proving — and, scored by how many of its problems our proofs covered at the time, that sample ranked first out of 2,000 random draws. Its high coverage measured our effort, not the truth. So we discarded it, fixed a fresh seed *before* drawing (recorded in the released `PREREGISTRATION.md`), and report whatever the new sample yields. On that sample, **71 of 100 problems** are proven to be one of the four ideas; the other 29 are the genuine tail. Of the 71, **69 are machine-checked here**; the remaining two—finding a minimum spanning tree (LC 1489) and finding the bridges of a graph (LC 1192)—rest on a classic algorithm whose correctness is itself a substantial formalization, so rather than redo it we cite an existing machine-checked proof [14, 16] that the algorithm is the graph computation we classify it as.

Checked against the editorials. To guard against modelling a strawman, every candidate was cross-checked against the problem’s published editorial. Three needed a note: two were re-labelled to a different one of the four (both still in scope), and Add Digits moved to the tail — its optimal accepted solution is a constant-time formula outside the four, so we count it conservatively

even though a single-pass solution also passes. That move is why the count is 71, not 72.

What this does and does not show. We prove that a solution of the right shape solves the problem; we do not prove it is byte-for-byte the platform’s accepted code (our model captures the solution’s structure, not its every line). So a certificate says “this problem really is solved by idea S ,” not “the accepted submission is literally this program.” How much each proof shows: **49 of the 69 prove the exact answer or a complete characterization** — the search finds a thing if and only if it is really there, an operation’s entire effect is captured in one equation, or the reported value is achieved by a real path or plan and beaten by none. The remaining 20 prove an exact law of the modelled solution; in each of those the model *is* the specification — the very recurrence, tally, or condition the accepted solution computes — so there is nothing further to prove about it. The submissions themselves are vouched for by the platform’s acceptance, which we never re-litigate. Two of the 20 (cycle detection; longest palindromic subsequence) additionally cite verified classical work for the underlying algorithm’s textbook correctness [15, 13]; unlike LC 1489 and LC 1192 above they *are* counted among the 69, since their own theorems are machine-checked here. This bar earns its keep: demanding the full balanced-string equivalence for bracket matching exposed a lenient model that silently accepted a stray closing bracket — the strict checker now carries the full equivalence, machine-checked. And we do not re-prove the deep theory behind the ideas—that shortest paths and dynamic programming are formally well-founded is already established [10, 11, 12, 13], which we cite rather than redo.

5 Related work

Two literatures bracket this paper. On the *empirical* side, software-engineering research has studied the interview *process* — developer sentiment about it [2], the effect of stress and format on performance [3], and how candidates prepare [4] — while preparation itself is organized by pattern catalogues [1] that enumerate surface families. We are not aware of prior work that tests whether those families are generatively distinct, or whether firms differ in the mix they ask. Closest in spirit, machine-learning work classifies competitive-programming problems by *predicted* solution type [5]; we settle the same kind of label by proof rather than prediction. On the *structural* side, the observation that wide families of programs are instances of a few recursion schemes is the Bird–Meertens tradition of program calculation [17] and its formal-methods descendants: associative aggregation over a sliding window is a verified fold [18], dynamic programming is verified memoized recursion [13], and shortest paths are a least fixpoint over a semiring [10, 11, 12]; verified-algorithms textbooks now develop these schemes end to end [19, 20]. The four-scheme decomposition itself is therefore not new—it is the classical Bird–Meertens repertoire. Our contribution is to connect that structural lens to the *empirical* distribution of interview problems, and to machine-check the scheme assignment per problem rather than assert it.

6 Threats to validity

One platform, truncated lists. The frequencies come from a single prep platform — we describe *commonly-asked* problems, not any firm’s private bar — and each firm’s list is its top-200 per recency window, so the rare tail is truncated and the company comparison runs over the heads of the distributions, where the data are densest. This aggregation also caps how different the firms *can* look: pooled, truncated, platform-curated lists may converge even if the firms’ true private bars differ, so the small between-firm V is best read as a property of the commonly-asked pool—possibly a side effect of the aggregation itself—not as strong evidence that the firms are interchangeable.

Borrowed labels. The surface-family tags are the platform’s—single-source and un-audited—so the “wide surface” count inherits that taxonomy’s splits. **We model the solution, not the submission.** As Section 4 notes, “this captures the accepted approach” is a modelling judgement — backed by the editorial cross-check, not certified by the proofs. **Which idea is a choice.** The proofs certify that an idea solves the problem (and, per `not_isFold_countDistinct`, that this is a claim a problem can provably fail), not that it is the only sensible description. **A single-sample estimate.** 71% comes from one pre-registered sample, so the 95% interval [61%, 79%], not the point value, is the honest summary; it errs low if anything, since the four ideas absorb some borderline “greedy”/“simulation” problems we left in the tail. **Window dependence.** The bias-corrected V ranges from below its per-window chance level in short windows to 0.125 at six months (still “small”), so “one shared syllabus” describes the pooled data, not every slice. **Two separate questions.** Coverage counts each distinct problem once; the company comparison weights by asking frequency. Both matter, and we keep them apart. **Shallow is not easy.** “Same idea” does not mean “same difficulty”—two single-pass problems can differ hugely in how hard the running value is to find. We claim shallow *structure*, not that the problems are easy.

7 Conclusion

Interview problems are wide on the surface but shallow underneath. On a pre-registered random sample, 71 of 100 problems (69 with a machine-checked proof here, 2 citing existing formal proofs) are one of four ideas (the single left-to-right pass is the most common); the rest are a genuine tail. Secondarily, in our pooled lists the seven firms’ mixes are barely distinguishable—though that similarity may owe something to the aggregation itself. For a candidate, the takeaway is blunt: learn the four ideas deeply rather than dozens of patterns shallowly; and, at least as far as these commonly-asked lists can show, preparing for “technical interviews” beats preparing for any one company.

Reproducibility and AI-use disclosure

All derived statistics, the pre-registration of the sampling seed (`PREREGISTRATION.md`, committed before the sample was drawn), the per-problem proof list from which the 71/100 re-derives, the editorial cross-check, and the complete Lean development are released at <https://github.com/corsur/swarm-tips/tree/main/research/wide-but-shallow>; the raw scrape alone is withheld under the platform’s terms. The proofs build with `lake build` with no gaps, `lake exe gate` recomputes every certificate verdict (Section 4), and `#print axioms` checks the standard-axioms claim. This paper was written by one human author with substantial AI assistance for drafting, statistical code, and the formalization; all claims, numbers, and proofs were checked by the author, and the machine-checked results are independently verifiable by re-running the proofs. We say so plainly: the paper’s central claims stand or fall on the released data and the proofs, not on how it was written.

References

- [1] S. Prashad. LeetCode Patterns. <https://seanprashad.com/leetcode-patterns/>, 2020.

- [2] M. Behroozi, S. Shirolkar, T. Barik, and C. Parnin. Hiring is broken: what do developers say about technical interviews? In *IEEE Symp. Visual Languages and Human-Centric Computing (VL/HCC)*, 2019.
- [3] M. Behroozi, S. Shirolkar, T. Barik, and C. Parnin. Does stress impact technical interview performance? In *Proc. ESEC/FSE*, 2020.
- [4] B. Bell, T. Thomas, S. W. Lee, and C. Brown. How do software engineering candidates prepare for technical interviews? *arXiv:2507.02068*, 2025.
- [5] R. C. A. Iacob et al. AlgoLabel: a large dataset for multi-label classification of algorithmic challenges. *Mathematics*, 8(11):1995, 2020.
- [6] H. Cramér. *Mathematical Methods of Statistics*. Princeton University Press, 1946.
- [7] W. Bergsma. A bias-correction for Cramér’s V and Tschuprow’s T . *Journal of the Korean Statistical Society*, 42(3):323–328, 2013.
- [8] L. de Moura and S. Ullrich. The Lean 4 theorem prover and programming language. In *Automated Deduction (CADE-28)*, 2021.
- [9] The mathlib Community. The Lean mathematical library. In *Proc. 9th ACM SIGPLAN Int. Conf. on Certified Programs and Proofs (CPP)*, 367–381, 2020.
- [10] B. Nordhoff and P. Lammich. Dijkstra’s shortest path algorithm. *Archive of Formal Proofs*, 2012.
- [11] A. Armstrong, G. Struth, and T. Weber. Kleene algebra. *Archive of Formal Proofs*, 2013.
- [12] D. J. Lehmann. Algebraic structures for transitive closure. *Theoretical Computer Science*, 4(1):59–76, 1977.
- [13] S. Wimmer, S. Hu, and T. Nipkow. Monadification, memoization and dynamic programming. *Archive of Formal Proofs*, 2018.
- [14] M. P. L. Haslbeck, P. Lammich, and J. Biendarra. Kruskal’s algorithm for minimum spanning forest. *Archive of Formal Proofs*, 2019.
- [15] P. Gammie. The Tortoise and the Hare algorithm. *Archive of Formal Proofs*, 2015.
- [16] P. Lammich and R. Neumann. A framework for verifying depth-first search algorithms. *Archive of Formal Proofs*, 2016; see also *Proc. CPP*, 137–146, 2015.
- [17] R. Bird and O. de Moor. *Algebra of Programming*. Prentice Hall, 1997.
- [18] L. Heimes, D. Traytel, and J. Schneider. Formalization of an algorithm for greedily computing associative aggregations on sliding windows. *Archive of Formal Proofs*, 2020.
- [19] T. Nipkow, J. Blanchette, M. Eberl, A. Gómez-Londoño, P. Lammich, C. Sternagel, S. Wimmer, and B. Zhan. *Functional Algorithms, Verified!* 2021. <https://functional-algorithms-verified.org>.
- [20] T. Nipkow, M. Eberl, and M. P. L. Haslbeck. Verified textbook algorithms: A biased survey. In *Automated Technology for Verification and Analysis (ATVA)*, 2020.

A Initial family-to-idea guess

The written-in-advance map from surface family to idea, by the shape of the canonical accepted solution: one pass carrying an accumulator $\rightarrow$ *streaming fold*; a recurrence over sub-results $\rightarrow$ *recursive decomposition (DP)*; iteration to a fixed point over a graph $\rightarrow$ *graph relaxation*; a monotone condition located by halving $\rightarrow$ *bisection*; anything else $\rightarrow$ *tail*. This is only the starting guess: the per-problem proof, checked against the editorial, settles each classification.